Process Scheduling

Purpose

The goal of this homework is to learn about process scheduling inside an operating system. You will work on the specified scheduling algorithm and simulate its performance.

Task

In this project, you will simulate the process scheduling part of an operating system. You will implement time-based scheduling, ignoring almost every other aspect of the OS. Please use message queues for synchronization.

Operating System Simulator

The operating system simulator, or OSS, will be your main program and serve as the master process. You will start the simulator (call the executable oss) as one main process who will fork multiple children at random times. The randomness will be simulated by a logical clock that will also be updated by oss.

In the beginning, oss will allocate shared memory for system data structures, including a process table with a process control block for each user process. The process control block is a fixed size structure and contains information to manage the child process scheduling. Notice that since it is a simulator, you will not need to allocate space to save the context of child processes. But you must allocate space for scheduling-related items such as total CPU time used, total time in the system, time used during the last burst, your local simulated pid, and process priority, if any. The process control block resides in shared memory and is accessible to the children. Since we are limiting ourselves to 20 processes in this class, you should allocate space for up to 18 process control blocks. Also create a bit vector, local to oss, that will help you keep track of the process control blocks (or process IDs) that are already taken.

oss simulates the passing of time by using a simulated system clock. The clock is stored in two shared integers in memory, one which stores seconds and the other nanoseconds. so use two unsigned integers for the clock. oss will be creating user processes at random intervals, say every second on an average. While the simulated clock (the two integers) is viewable by the child processes, it should only be advanced (changed) by oss.

Note that oss simulates time passing in the system by adding time to the clock and as it is the only process that would change the clock, if a user process uses some time, oss should indicate this by advancing the clock. In the same fashion, if oss does something that should take some time if it was a real operating system, it should increment the clock by a small amount to indicate the time it spent.

oss will create user processes at random intervals (of simulated time), so you will have two constants; let us call them `maxTimeBetweenNewProcsNS` and `maxTimeBetweenNewProcsSecs`. oss will launch a new user process based on a random time interval from 0 to those constants. It *generates* a new process by allocating and initializing the process control block for the process and then, `forks` the process. The child process will `execl` the binary. I would suggest setting these constants initially to spawning a new process about every 1 second, but you can experiment with this later to keep the system busy. New processes that are created can have one of two scheduling classes, either real-time or a normal user process, and will remain in that class for their lifetime. There should be constant representing the percentage of time a process is launched as a normal user process or a real-time one. While this constant is specified by you, it should be heavily weighted to generating mainly user processes.

oss will run concurrently with all the user processes. After it sets up all the data structures, it enters a loop where it generates and schedules processes. It *generates* a new process by allocating and initializing the process control block for the process and then, `forks` the process. The child process will `execl` the binary.

oss will be in control of all concurrency. In the beginning, there will be no processes in the system but it will have a time in the future where it will launch a process. If there are no processes currently ready to run in the system, it should increment the clock until it is the time where it should launch a process. It should then set up that process, generate a new time where it will launch a process and then using a message queue, schedule a process to run by sending it a message. It should then wait for a message back from that process

that it has finished its task. If your process table is already full when you go to generate a process, just skip that generation, but do determine another time in the future to try and generate a new process.

Advance the logical clock by 1.xx seconds in each iteration of the loop where xx is the number of nanoseconds. xx will be a random number in the interval [0,1000] to simulate some overhead activity for each iteration.

A new process should be generated every 1 second, on an average. So, you should generate a random number between 0 and 2 assigning it to time to create new process. If your clock has passed this time since the creation of last process, generate a new process (and `execl` it).

`oss` acts as the scheduler and so will *schedule* a process by sending it a message using a message queue. When initially started, there will be no processes in the system but it will have a time in the future where it will launch a process. If there are no processes currently ready to run in the system, it should increment the clock until it is the time where it should launch a process. It should then set up that process, generate a new time where it will create a new process and then using a message queue, schedule a process to run by sending it a message. It should then wait for a message back from that process that it has finished its task. If your process table is already full when you go to generate a process, just skip that generation, but do determine another time in the future to try and generate a new process.

Scheduling Algorithm

Assuming you have more than one process in your simulated system, `oss` will *select* a process to run and *schedule* it for execution. It will select the process by using a scheduling algorithm with the following features:

Implement a multi-level feedback queue with Linux-like scheduling. That is, you will have two classes of processes: active and expired. When a process has exhausted its quantum, it will be moved to expired set. `oss` will always pick up the process with highest priority that is acive. There are four scheduling queues, each having an associated time quantum. The base time quantum is determined by you as a constant, let us say something like 10 milliseconds, but certainly could be experimented with. The highest priority queue has this base time quantum as the amount of time it would schedule a child process if it scheduled it out of that queue. The second highest priority queue would have half of that, the third highest priority queue would have quarter of the base queue quantum and so on, as per a normal multi-level feedback queue. If a process finishes using its entire timeslice for the queue, it should be moved to a queue one lower in priority. If a process comes out of a blocked queue, it should go to the highest priority queue.

When `oss` has to pick a process to schedule, it will look for the highest priority occupied queue and schedule the process at the head of this queue. The process will be *dispatched* by sending the process a message using a message queue indicating how much of a time slice it has to run. Note that this scheduling itself takes time, so before launching the process the `oss` should increment the clock for the amount of work that it did, let us say from 100 to 10000 nanoseconds.

User Processes

All user processes are alike but simulate the system by performing some tasks at random times. The user process will keep checking in the shared memory location if it has been scheduled and once scheduled, it will start to run. It should generate a random number to check whether it will use the entire quantum, or only a part of it (a binary random number will be sufficient for this purpose). If it has to use only a part of the quantum, it will generate a random number in the range [0,quantum] to see how long it runs.

The user processes will wait on receiving a message giving them a time slice and then it will simulate running. They do not do any *actual* work but instead send a message to `oss` saying how much time they used and if they had to use I/O or had to terminate.

As a constant in your system, you should have a probability that a process will terminate when scheduled. I would suggest this probability be fairly small to start. Processes will then, using a random number, use this to determine if they should terminate. Note that when creating this random number you must be careful that the seed you use is different for all processes, so I suggest seeding off of some function of a processes pid. If it would terminate, it would of course use some random amount of its timeslice before terminating. It should indicate to `oss` that it has decided to terminate and also how much of its timeslice it used, so that `oss` can increment the clock by the appropriate amount.

Once it has decided that it will not terminate, then we have to determine if it will use its entire timeslice or if it will get blocked on an event. This should be determined by a random number between 0 and 1. If it uses up its timeslice, this information should be conveyed to master. Otherwise, the process starts to wait for an event that will last for $r.s$ seconds where r and s are random numbers with range $[0, 3]$ and $[0, 1000]$ respectively, and 3 indicates that the process gets preempted after using $p\%$ of its assigned quantum, where p is a random number in the range $[1, 99]$. As this could happen for multiple processes, this will require a blocked queue, checked by `oss` every time it makes a decision on scheduling to see if it should wake up these processes and put them back in the appropriate

queues. Note that the simulated work of moving a process from a blocked queue to a ready queue would take more time than a normal scheduling decision so it would make sense to increment the system clock to indicate this.

Your simulation should end with a report on average wait time, average CPU utilization and average time a process waited in a blocked queue. Also include how long the CPU was idle with no ready processes.

Make sure that you have signal handing to terminate all processes, if needed. In case of abnormal termination, make sure to remove shared memory and message queues.

Invoking the solution

Execute oss with no parameters. You may add some parameters to modify the number of processes or base quantum but if you do so, document this in your README.

Log Output

Your program should send enough output to a log file such that it is possible for me to determine its operation. For example:

```
OSS: Generating process with PID 3 and putting it in queue 1 at time 0:5000015
OSS: Dispatching process with PID 2 from queue 1 at time 0:5000805,
OSS: total time this dispatch was 790 nanoseconds
OSS: Receiving that process with PID 2 ran for 400000 nanoseconds
OSS: Putting process with PID 2 into queue 2
OSS: Dispatching process with PID 3 from queue 1 at time 0:5401805,
OSS: total time this dispatch was 1000 nanoseconds
OSS: Receiving that process with PID 3 ran for 270000 nanoseconds,
OSS: not using its entire time quantum
OSS: Putting process with PID 3 into queue 1
OSS: Dispatching process with PID 1 from queue 1 at time 0:5402505,
OSS: total time spent in dispatch was 7000 nanoseconds
etc
```

I suggest not simply appending to previous logs, but start a new file each time. Also be careful about infinite loops that could generate excessively long log files. So for example, keep track of total lines in the log file and terminate writing to the log if it exceeds 10000 lines.

Note that the above log was using arbitrary numbers, so your times spent in dispatch could be quite different.

Suggested implementation steps

I highly suggest you do this project incrementally. I suggest implementation in the following order:

- Start by creating a Makefile that compiles and builds the two executables: oss and user_proc.

- Have oss create a process control table with one user process (of real-time class) to verify it is working

- Schedule the one user process over and over, logging the data

- Create the round robin queue, add additional user processes, making all user processes alternate in it

- Keep track of and output statistics like throughput, idle time, etc

- Implement an additional user class and the multi-level feedback queue.

- Add the chance for user processes to be blocked on an event, keep track of statistics on this

Do not try to do everything at once and be stuck with no idea what is failing.

Termination Criteria

`oss` should stop generating processes if it has already generated 100 processes or if more than 3 real-life seconds have passed. If you stop adding new processes, the system should eventually empty of processes and then it should terminate. What is important is that you tune your parameters so that the system has processes in all the queues at some point and that I can see that in the log file. As discussed previously, ensure that appropriate statistics are displayed.

Criteria for success

Make sure that you implement the specified algorithm and document it appropriately. *You must clean up after yourself*. That is, after the program terminates, whether normally or by force, there should be no shared memory, semaphore, or message queue that is left allocated to you.

Grading

1. *Overall submission: 30pts*. Program compiles and upon reading, seems to solve the assigned problem in the specified manner.

2. *README/Makefile: 10pts*. Ensure that they are present and work appropriately.

3. *Code readability: 10pts*. Code should be readable with appropriate comments. Author and date should be identified.

4. *Conformance to specifications: 50pts*. Algorithm is properly implemented and documented.

Submission

Handin an electronic copy of all the sources, `README`, Makefile(s), and results. Create your programs in a directory called *username*.4 where *username* is your login name on hoare. Once you are done with everything, *remove the executables and object files*, and issue the following commands:

```
% cd
% chmod 755 ~
% ~sanjiv/bin/handin cs4760 4
% chmod 700 ~
```

Do not forget `Makefile` (with suffix rules), version control, and `README` for the assignment.